PDF Answering AI

Problem Statement and Objective

Problem Statement

The task is to develop an AI model capable of answering questions based on the content of a PDF document. This model should leverage natural language processing (NLP) techniques to understand the context and provide accurate responses to queries.

Objective

The objective of this project is to create a system that can interactively respond to questions about a PDF document, mimicking a scenario where the PDF "talks" back to the user. This involves fine-tuning a pre-trained model (the model is bert-large-uncased-whole-word-masking) on the SQuAD (Stanford Question Answering Dataset) dataset to ensure the model understands and processes the context of the text within the PDF to generate precise answers.

Methodology

1. **Data Preparation**:
    - Collected relevant data by using the SQuAD dataset, which is a large dataset for training question-answering systems.
    - Pre-processed the data to ensure compatibility with the chosen model architecture.
2. **Model Selection and Fine-Tuning**:
    - Selected a pre-trained model, BERT (Bidirectional Encoder Representations from Transformers), due to its effectiveness in NLP tasks. [The specific model is bert-large-uncased-whole-word-masking].
    - Fine-tuned BERT on the SQuAD dataset to adapt it to the task of answering questions based on the context of the PDF content.
3. **Implementation**:
    - Developed a nlp pipeline from transformers to upload PDFs and extract text.
    - Implemented a mechanism to input questions and retrieve relevant text sections using the fine-tuned model.
    - Developed the frontend of the project using streamlit making it easy to interact.

Failed Approaches

1. **Vector Databases and Cosine Similarity**:
    - Attempted to convert PDF text into vectors using word2vec and GloVe embeddings.
    - Transformed questions into vector format and tried to find the closest match using cosine similarity.
    - This approach failed due to a lack of contextual understanding and difficulty in handling complex queries.
2. **Finding the correct pre-trained model and the best dataset for fine tuning:**
    - Some models were not up to the mark when it came to predicting answers.
    - Also in the fine tuning some datasets were very small and even fine tuning on them took a lot of time and yet they were not able to generate the correct or the expected answers.
3. **Other Attempts**:
    - **Using TF-IDF**: Tried using term frequency-inverse document frequency (TF-IDF) to identify important words and match queries, but it did not capture the context effectively.
    - **LSTM Networks**: Experimented with Long Short-Term Memory (LSTM) networks, but they struggled with the large size of text and required extensive computational resources.

Results:

I answered the questions from Aries Open Projects ML.pdf itself and below are some examples.

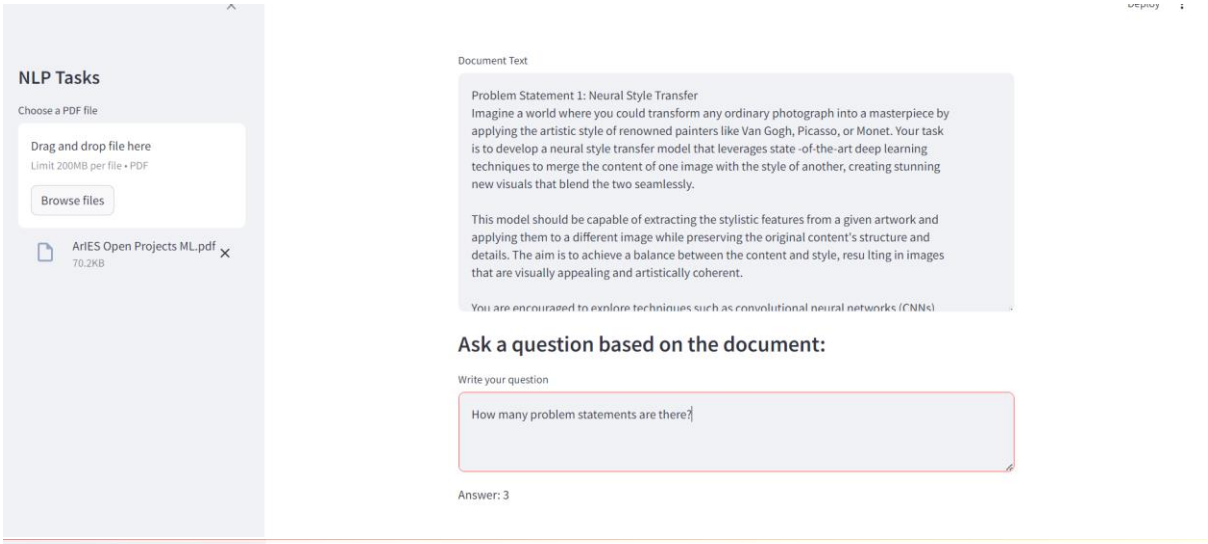

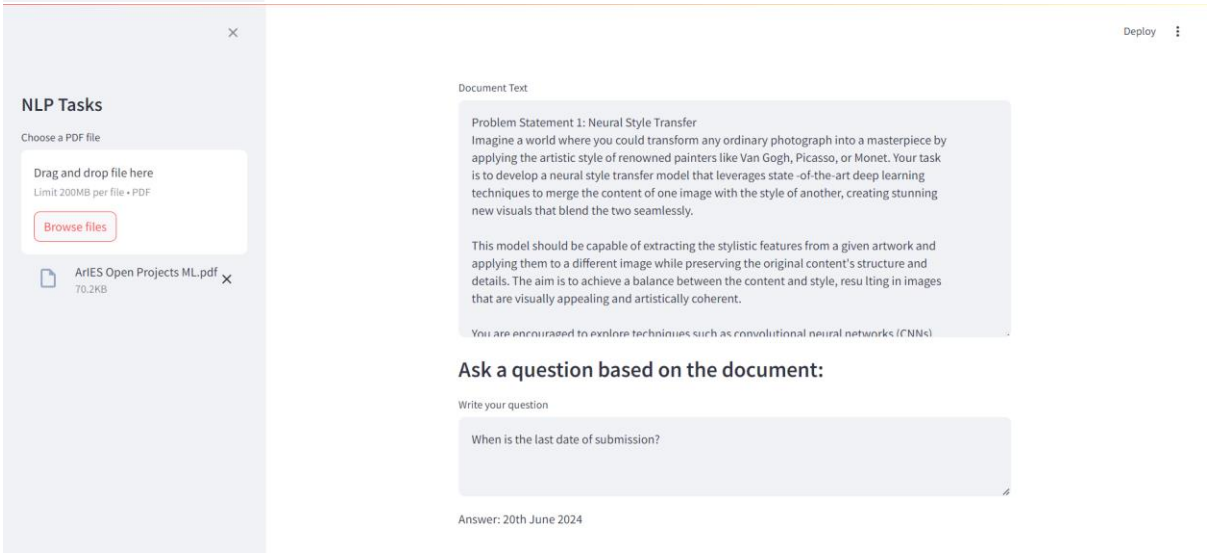

Also here are some results from training the model which I get from **“wandb”** which helps in visualising the training phase of the model and the system processing during that phase.

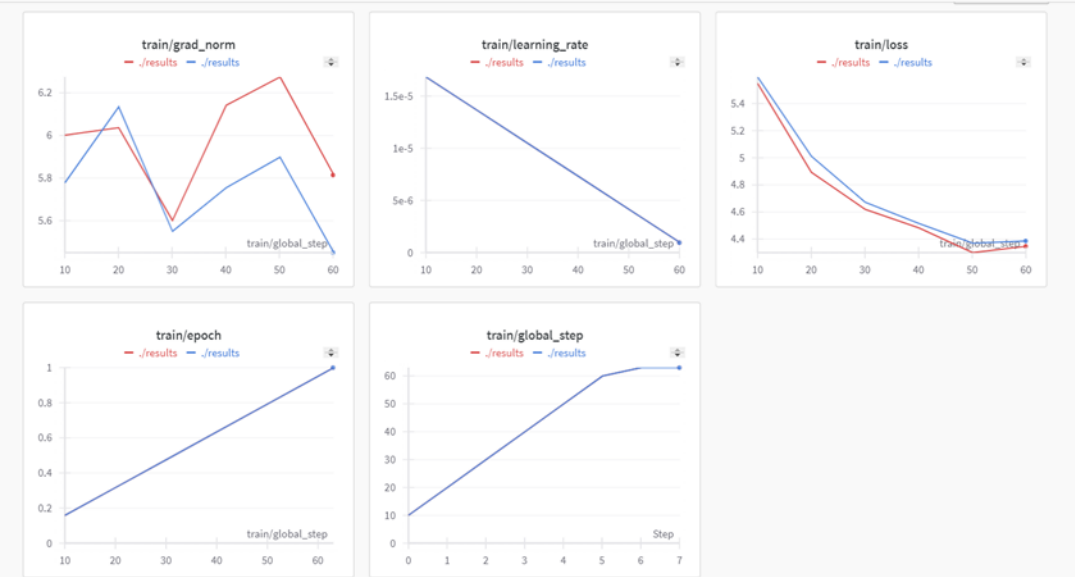

System 11

Add panel

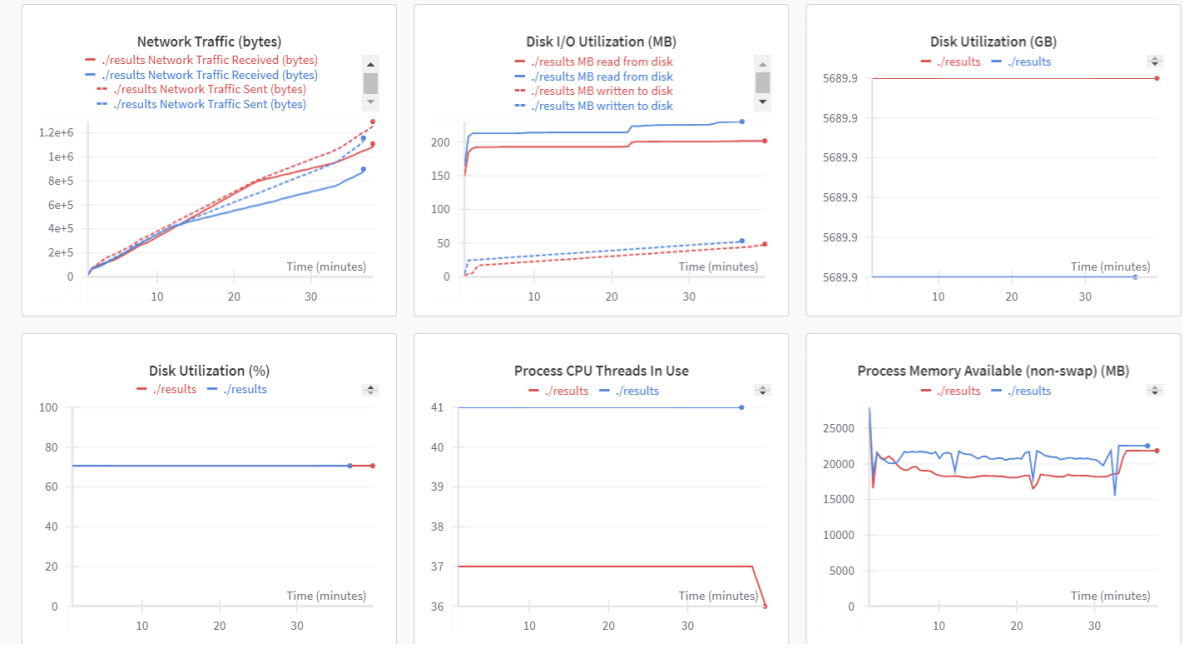

Analysis and Insights

Learnings and Insights

1. **Understanding of Transformers**:
    - Gained a deep understanding of transformer models, particularly BERT, and their application in NLP tasks.
    - Learned about the architecture and functioning of transformers, including attention mechanisms and their advantages over traditional RNNs and LSTMs.
2. **Applications and Uses**:
    - Realized the broad applicability of transformer models in various NLP tasks beyond question answering, such as text summarization, translation, and sentiment analysis.
    - Understood the importance of fine-tuning pre-trained models on domain-specific data to improve performance.
3. **Challenges and Solutions**:
    - Encountered challenges with handling large volumes of text and ensuring the model's responses were accurate and contextually relevant.
    - Overcame these challenges by fine-tuning on a relevant dataset and optimizing the text extraction and processing pipeline.
    - Also, a major challenge was the time taken for fine tuning the model. SQuAD is a large dataset and fine tuning a large model like bert-large-uncased-whole-word-masking takes lot of time and patience.

Summary and Future Improvements

Summary

In this project, we developed an AI model that can answer questions based on the content of a PDF document. By fine-tuning a pre-trained BERT model on the SQuAD dataset, we achieved a system capable of understanding and responding to queries accurately. This project provided valuable insights into the workings of transformer models and their practical applications in NLP tasks.

Future Improvements

1. **Enhanced Contextual Understanding**:
    - Integrate more sophisticated text preprocessing techniques to improve the model's understanding of the document structure.
    - Experiment with other transformer models like RoBERTa for potentially better performance.
2. **User Interface Improvements**:
    - Develop a more user-friendly interface for interacting with the model, such as a browser extension or a mobile app.
    - Implement voice input for questions to make the system more accessible.
3. **Expand Dataset**:
    - Fine-tune the model on additional datasets to cover a wider range of topics and improve generalization.

References

1. **SQuAD Dataset**: https://rajpurkar.github.io/SQuAD-explorer/
2. **Attention is All You Need**: https://arxiv.org/abs/1706.03762
3. **BERT Playlist:** https://www.youtube.com/playlist?list=PLam9sigHPGwOBuH4_4fr-XvDbe5uneaf6
4. **Sequence model by Andrew Ng:** https://www.youtube.com/watch?v=S7oA5C43Rbc&t=5130s
5. **Illustrated Transformer**: http://jalammar.github.io/illustrated-transformer/

REPORT BY:

RAGHAV NYATI
22112082